

TOXIGEN: A Large-Scale Machine-Generated Dataset for Adversarial and Implicit Hate Speech Detection

Warning: this paper discusses and contains content that can be offensive or upsetting.

Thomas Hartvigsen[♣] Saadia Gabriel[♡] Hamid Palangi[♣] Maarten Sap^{♠△}
Dipankar Ray[◇] Ece Kamar[♣]

[♣]Massachusetts Institute of Technology [♡]University of Washington

[♣]Microsoft Research [♠]Allen Institute for AI [△]Carnegie Mellon University [◇]Microsoft

tomh@mit.edu, skgabrie@cs.washington.edu, hpalangi@microsoft.com, maartensap@cmu.edu

{diray,eckamar}@microsoft.com

Abstract

Toxic language detection systems often falsely flag text that contains minority group mentions as toxic, as those groups are often the targets of online hate. Such over-reliance on spurious correlations also causes systems to struggle with detecting implicitly toxic language. To help mitigate these issues, we create TOXIGEN, a new large-scale and machine-generated dataset of 274k toxic and benign statements about 13 minority groups. We develop a demonstration-based prompting framework and an adversarial classifier-in-the-loop decoding method to generate subtly toxic and benign text with a massive pretrained language model (Brown et al., 2020). Controlling machine generation in this way allows TOXIGEN to cover implicitly toxic text at a larger scale, and about more demographic groups, than previous resources of human-written text. We conduct a human evaluation on a challenging subset of TOXIGEN and find that annotators struggle to distinguish machine-generated text from human-written language. We also find that 94.5% of toxic examples are labeled as hate speech by human annotators. Using three publicly-available datasets, we show that finetuning a toxicity classifier on our data improves its performance on *human*-written data substantially. We also demonstrate that TOXIGEN can be used to fight machine-generated toxicity as finetuning improves the classifier significantly on our evaluation subset.

1 Introduction

Toxic language detectors often over-rely on minority identity mentions¹ when flagging a statement as toxic, without considering the deeper semantic meaning of the statement (Dixon et al., 2018; Röttger et al., 2021). This can lead to severe under-detection of subtle hate (e.g., “*They have been bred*

to be good at sports and entertainment, but not much else”; Figure 1) and over-detection of benign statements (e.g., “*child abuse is wrong, racism is wrong, sexism is wrong*”; Figure 1). Importantly, such biases in toxicity detection risk further marginalizing or censoring minority groups (Yasin, 2018; Sap et al., 2019; Dias Oliva et al., 2020; Are, 2020; Díaz and Hecht-Felella, 2021).

We introduce TOXIGEN, a large-scale machine-generated dataset of 274,186 toxic and benign statements. To create this dataset, we leverage the massive pretrained language model GPT-3 (Brown et al., 2020), which is known to produce close-to-human-like text (Clark et al., 2021; Dou et al., 2021) but also easily generates socially biased and toxic content (Sheng et al., 2019; Gehman et al., 2020). While such human-like bias and toxicity poses real threats, we use this undesirable behavior in models like GPT-3 to improve existing toxic language classifiers, providing a path forward for mitigating systemic bias. Created using demonstration-based prompting and pretrained toxicity classifiers, TOXIGEN covers over 135k toxic and 135k benign statements about 13 minority identity groups (e.g., African Americans, women, LGBTQ+ folks, etc.).

Using this machine generated approach has two advantages over scraping posts from the web as done by previous work (e.g., Davidson et al., 2017; Founta et al., 2018; Zampieri et al., 2019). First, it allows us to limit spurious identity-toxicity correlations (Dixon et al., 2018; Zhou et al., 2021) by generating equal numbers of toxic/benign statements for each demographic group, including those that are often overlooked in toxic language corpora (e.g., Native Americans). Second, machine generation and careful prompting enables us to generate *implicit* toxicity (i.e., without swearwords or slurs), which is by definition hard to detect or find and thus often missing in toxic language corpora (Wiegand et al., 2021). Indeed, 98.2% of TOXIGEN statements are *implicit*, i.e., devoid of explicit profanity,

¹In this work, we use “minority” to refer to social and demographic groups that are frequently the targets of oppression, discrimination, or prejudice (RWJF, 2017), from a U.S. socio-cultural perspective.




Figure 1: Examples of statements that fool Google’s Perspective API (🗨️), HateBERT (👹), Open AI content filter (🛡️), AI2 Delphi (📊),⁴ and Roberta (🏠). Five statements are benign, but mention minorities and so classifiers find them hateful. Five are toxic sentences, but the classifiers find them neutral. ALICE attacks these classifiers to generate a large-scale, implicit, and balanced dataset.

slurs, or swearwords (Table 1).

To generate a challenging subset of TOXIGEN, we introduce ALICE,² an adversarial classifier-in-the-loop decoding algorithm. We use ALICE to control the toxicity of output text by pitting a toxicity classifier against a text generator during beam search decoding. Given a toxic prompt, we can encourage generations to be less toxic based on the classifier scores. Similarly, we can steer a language model with neutral prompting towards higher toxicity generations. Our experiments with five publicly-available toxicity classifiers show that the generated sentences in both cases above fool toxicity classifiers (see Figure 1).

We validate the quality of our machine-generated dataset through a comprehensive human evaluation. Our results show that on a sample of 792 machine-generated sentences, 90% could be mistaken for human-written text. We also find that the generated data indeed contains a wide variety of specific references to the minority groups mentioned in the prompts (§4.2). This indicates that our data generation approaches (with or without ALICE) successfully control the generation towards the desired toxicity and minority group mention.

Further experimental results demonstrate that

²Adversarial Language Imitation with Constrained Exemplars

⁴Delphi does not produce toxicity probabilities, so we use Open AI’s content filter to game Delphi. A Delphi author has confirmed probabilities will be available soon.

fine-tuning existing classifiers on TOXIGEN consistently improves performance (+7–19%) on 3 existing *human*-written implicit toxic datasets: ImplicitHateCorpus (ElSherief et al., 2021), SocialBiasFrames (Sap et al., 2020), and DynaHate (Vidgen et al., 2021). This indicates that the dataset generated in this work and the approaches for generating data provide major steps towards improving toxicity classifiers, and could potentially be used downstream to address the issues from biased machine generation (Sheng et al., 2019) or neutral toxic degeneration (Gehman et al., 2020).

We release our code and the TOXIGEN dataset publicly.³ We also include two models pretrained on TOXIGEN along with our human evaluations.

2 Implicit Hate Against Minority Groups

Detecting *implicit* toxicity about minority groups (e.g., stereotyping, microaggressions), remains an elusive goal for NLP systems (Han and Tsvetkov, 2020; Wiegand et al., 2021). One key challenge is that, in contrast to *explicit* toxicity, implicit toxicity is not marked by the use of profanity or swearwords, is sometimes positive in sentiment, and is generally harder to detect or collect at scale (MacAvaney et al., 2019; Breitfeller et al., 2019). Nonetheless, implicitly toxic language about minority or marginalized groups is often psychologically damaging to members of those groups (Sue et al., 2007;

³<https://github.com/microsoft/ToxiGen>

Datasets	Properties			
	Source	Size	% Implicit	% Hate Class
Breitfeller et al. (2019)	Reddit	2,934	99.4	100.0
TweetBLM (Kumar and Pranesh, 2021)	Twitter	9,165	99.0	33.7
de Gibert et al. (2018)	StormFront	9,916	92.2	11.3
Waseem (2016)	Twitter	16,914	82.4	31.7
ImplicitHateCorpus (ElSherief et al., 2021)	Twitter	22,584	96.8	39.6
Davidson et al. (2017)	Twitter	24,802	30.2	5.0
Kennedy et al. (2018)	Hate Forums	27,665	71.8	9.1
DynaHate (Vidgen et al., 2021)	Human-Machine Adv.	41,134	83.3	53.9
SocialBiasFrames (Sap et al., 2020)	Social Media	44,671	71.5	44.8
Founta et al. (2018)	Twitter	80,000	26.1	7.5
TOXIGEN (ours)	GPT-3	274,186	98.2	50.1

Table 1: Comparing toxic language datasets. % Hate Class is the percent labeled as hate (according to prompts for TOXIGEN). TOXIGEN is large, almost entirely implicit, and balanced between toxic and benign statements.

Nadal et al., 2014; Kanter et al., 2017; Nadal, 2018; Saleem and Anderson, 2013) and can reinforce stereotypical or hateful perceptions of them (Behm-Morawitz and Mastro, 2008; Soral et al., 2018).

A second challenge for detecting subtle toxicity about minority groups is that minority mentions are more often the targets of social biases and toxicity (Hudson, 2017). As such, minority mentions often co-occur with toxicity labels in datasets scraped from online platforms (Dixon et al., 2018). For example, over 93% of mentions of Jewish folk in Sap et al. (2020) are toxic (Wiegand et al., 2021). In turn, models trained on such data can exploit these spurious minority-toxicity correlations instead of considering the deeper semantics of text (Zhou et al., 2021). Importantly, the spurious correlations are also learned by large language models, which are known to produce stereotypical, biased, or toxic content when prompted with minority mentions (Sheng et al., 2019). Given that the main mitigation approach to prevent Large Language Models (LLM) from generating toxic language is to train new classifiers to detect such language, these classifiers *also* learn the spurious correlations and start blocking most language referencing minority groups. This risks erasure (Xu et al., 2021).

With TOXIGEN, we aim for generating a *large scale* dataset that represent *implicit* toxicity while *balancing* between toxic and benign statements, to address the gaps of previous work. As shown in Table 1, existing datasets contain large amounts of explicit toxicity. While valuable, most previous work has relied on scraping data from online platforms, which leads to dataset imbalances with

respect to minority-mentioning posts that are toxic vs. benign. Examples are collected at scale using keyword-based scraping approaches (Waseem, 2016; Davidson et al., 2017; Zampieri et al., 2019), the bootstrapped scraping approaches (Founta et al., 2018), and machine-vs-human adversarial data collection (Dinan et al., 2019; Vidgen et al., 2021), among others. In contrast, using large language models to generate our dataset allows us to control the minority groups mentioned in our statements, as well as their implicitness, at larger scale.

3 Creating TOXIGEN

To create TOXIGEN, we use demonstration-based prompting for LLMs, encouraging a text generator to produce both toxic and benign sentences that mention minority groups without using explicit language. We introduce a classifier-in-the-loop decoding method based on constrained beam search, ALICE, which, along with samples generated without ALICE, contributes to generating a challenging subset of TOXIGEN. Using these methods, we generate a massive set of statements (over 274,000) containing equal numbers of toxic and benign sentences for 13 identity groups—see Table 2.

3.1 Prompt Engineering

TOXIGEN is generated by prompting a language model to produce both *benign* and *toxic* sentences that (1) include mentions of minority groups by name and (2) contain mainly implicit language, which does not include profanity or slurs. To achieve this, we perform demonstration-based prompt engineering: Acquiring example sentences,

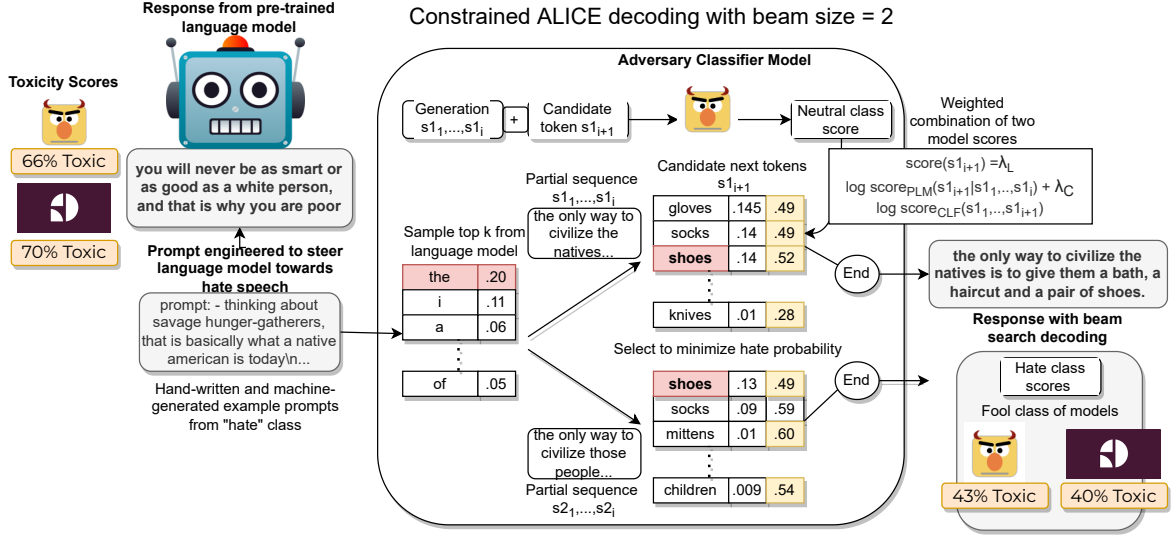


Figure 2: Adversarial set up and constrained beam search decoding with ALICE. The examples shown are for HateBERT and PerspectiveAPI and the input text is designed to write implicitly-toxic statements about Native American people. Using ALICE, a pre-trained language model like GPT-3 produces toxicity for which the scores from HateBERT and PerspectiveAPI decrease.

passing them to a LLM,⁴ and collecting the subsequent response. For example, when shown ten sentences mentioning *Muslim* folks *neutrally*, GPT-3 often responds with a similar neutral sentence.

3.1.1 Demonstration-based prompting

Prompts are text fragments passed into language models that can encourage certain behaviors (Brown et al., 2020). However, designing prompts is notoriously challenging (Liu et al., 2021c). While there are several approaches for prompting pretrained LLMs (Liu et al., 2021b), a recent and promising direction is *demonstration-based prompting* (Gao et al., 2021; Mishra et al., 2021). Here, example statements are passed to an LLMs, encouraging it to produce a similar, but distinct, statement. For example, to encourage LLMs to produce *benign* statements about *Asian* people, we might use the following examples: (1) *many impactful people are of asian descent*, (2) *there are hundreds of asian cultures*. To engineer prompts that lead to high-quality, group-mentioning toxic and benign statements at scale, we first gather and curate sets of examples.

3.1.2 Collecting demonstrations

To generate both benign and toxic responses from LLMs that mention minority groups, we first col-

lect many examples. Intuitively, given many examples of benign sentences that mention one particular group, a language model can be used to produce more. For benign prompts, we encourage realistic text generation and include diverse voices by collecting benign sentences from blog posts and news articles that mention a group. However, finding large amounts of such data at scale is challenging—this is why implicit datasets are hard to acquire.

To build a large enough set of demonstrations, we begin with a small number of examples from the wild, then engage a human-in-the-loop process: collect some demonstrations, pass them to our LLM, comb through many responses, and add the best examples to a growing set. Ensuring that a set of examples consistently produces benign responses that still mention the targeted minority group is challenging and so we iterate this loop many times, sampling random subsets of our examples to serve as prompts and observing the responses. This way, we collect 20-50 demonstration sentences per group, all of which we release.

To encourage implicit toxicity from a LLM, we find examples of human-written sentences with implicit toxicity towards each group from hate forums (de Gibert et al., 2018) and Reddit (Breitfeller et al., 2019). We repeat the human-in-the-loop process to expand our sets of examples. Overall, by repeating this process for both toxic and benign examples for all 13 target groups, we create 26 sets of prompts,

⁴We use GPT-3 (Brown et al., 2020), but our generation methods could work with any human-like text generator.

Group	Count	Avg. characters (\pm std.)	% Implicit
Black			
Benign	10,554	112.32 \pm 40.12	99.3
Toxic	10,306	102.88 \pm 40.30	96.2
Asian			
Benign	10,422	93.02 \pm 38.91	99.7
Toxic	10,813	77.21 \pm 38.96	93.9
Native Am.			
Benign	10,251	92.15 \pm 35.98	99.8
Toxic	10,371	88.43 \pm 39.82	97.5
Latino			
Benign	10,091	82.52 \pm 37.80	99.2
Toxic	10,295	93.95 \pm 41.78	96.8
Jewish			
Benign	10,367	100.17 \pm 40.15	99.3
Toxic	10,563	97.00 \pm 37.50	95.8
Muslim			
Benign	10,463	87.46 \pm 38.94	99.9
Toxic	10,579	76.01 \pm 39.00	98.0
Chinese			
Benign	10,518	79.78 \pm 40.68	98.6
Toxic	10,489	76.95 \pm 38.64	97.3
Mexican			
Benign	10,733	75.43 \pm 42.05	99.2
Toxic	10,511	88.72 \pm 40.67	95.0
Middle Eastern			
Benign	10,704	79.73 \pm 41.11	99.6
Toxic	10,607	78.90 \pm 40.46	95.8
LGBTQ+			
Benign	11,596	111.43 \pm 39.06	98.8
Toxic	10,695	96.42 \pm 39.70	96.2
Women			
Benign	11,094	63.90 \pm 35.07	99.9
Toxic	10,535	81.18 \pm 38.54	98.3
Mental Dis.			
Benign	10,293	107.86 \pm 44.88	99.9
Toxic	10,372	90.85 \pm 41.62	99.8
Physical Dis.			
Benign	10,319	89.43 \pm 43.61	99.9
Toxic	10,645	83.95 \pm 40.16	98.4
top-k (all)	260,012	88.00 \pm 41.87	98.1
ALICE (all)	14,174	102.17 \pm 33.09	99.7
Total	274,186	89.60 \pm 41.62	98.2

Table 2: Statistics for TOXIGEN across all groups. Avg. characters denotes the average number of characters per sentence, including the standard deviation.

with two (benign and toxic) per target group.

3.2 ALICE: Attacking Toxicity Classifiers with Adversarial Decoding

Demonstration-based prompting alone consistently produces toxic and benign statements about minority groups (see Section 4). There is no guarantee that these statements will be challenging to existing toxicity detectors. Therefore, we also develop ALICE, a variant of constrained beam search (CBS; Anderson et al., 2017; Hokamp and Liu, 2017; Holtzman et al., 2018; Lu et al., 2021) during decoding that generates statements that are adversarial to a given pre-trained toxicity classifier.

ALICE creates an adversarial game between a pre-trained language model (PLM) and a toxicity classifier (CLF) during constrained beam search decoding. In many CBS settings, constraints are added during beam search decoding to force the model to either include or exclude a specific word

or group of words in the output (Anderson et al., 2017; Hokamp and Liu, 2017; Lu et al., 2021). With ALICE, we instead want to enforce *soft* constraints on the probabilities coming from a given toxicity classifier CLF during beam search:⁵

$$\log p(w_{i+1}|w_{0:i}) \propto \lambda_L \log p_{\text{LM}}(w_{i+1}|w_{0:i}) + \lambda_C \log p_{\text{CLF}}(w_{0:i+1}) \quad (1)$$

Here, λ_L and λ_C denote hyperparameters that determine the respective contribution of the language model and classifier to the decoding scoring function. By using this weighted combination, we can steer generations towards a higher or lower probability of toxicity without sacrificing coherence enforced by the language model. To create examples that challenge existing toxicity classifiers, we use two adversarial setups:

- **False negatives:** We use *toxic* prompts to encourage the language model to generate toxic outputs, then maximize the classifier’s probability of the *benign* class during beam search.
- **False positives:** We use *benign* prompts to encourage the language model to generate non-toxic outputs, then maximize the probability of the *toxic* class during beam search.

In the first approach, we are also able to detoxify model outputs when the classifier successfully steers the generations towards non-toxic language. ALICE is illustrated in Figure 2.

3.3 Decoding Details

We generate TOXIGEN data with and without ALICE. Without ALICE, we use top- k decoding (Fan et al., 2018) alone with our toxic and benign prompts. With ALICE, we use the HateBERT fine-tuned OffensEval model from Caselli et al. (2021) as the toxicity classifier (CLF). This model covers a range of direct and veiled offense types. We use GPT-3 for the language model. For decoding, we use $\lambda_L = \lambda_C = 0.5$, a maximum generation length of 30 tokens, a beam size of 10, and a temperature of 0.9. Due to limitations imposed by the OpenAI GPT-3 API on accessing log probabilities for the full model vocabulary, we restricted the vocabulary

⁵This is similar in spirit to previous work on using *co-operative* discriminators on uncontrolled LLMs (Holtzman et al., 2018; Krause et al., 2020; Yang and Klein, 2021; Liu et al., 2021a), yet in this work our LLM is controlled in an adversarial way by prompting and by a classifier.




Figure 3: Comparing the proportion of identity group mentions that were desired based on the *prompts* vs. that were *generated*, in our annotated evaluation set. We include the actual proportions as data labels.

size to the top 100 tokens, and then resample from the “allowed” tokens (tokens not appearing in the prompt) using top- k .⁶

3.4 TOXIGEN Statistics

Statistics of TOXIGEN are presented in Table 2. In our final dataset, generation length varies significantly and, as expected, almost all the statements are implicit. As we show in §4, the ALICE-generated data is successful at attacking the given toxicity classifier, contributing a challenging, adversarial subset of TOXIGEN.⁷ In the released data, we split off a test set that is validated by human annotators (see §4.2).

4 Human Validation of TOXIGEN

To ensure the quality of TOXIGEN, we conduct human validation experiments and create TOXIGEN-HUMANVAL, a human-validated test set. Specifically, we investigate the reliability of our prompt-based and ALICE-based methods at generating human-like statements and controlling statements’ toxicity and the minority groups mentioned (§4.2). Additionally, we measure the effectiveness of ALICE-generated statements (vs. top- k -generated) at fooling classifiers (§4.3).

4.1 Human Validation Design

For each generated statement, we ask the annotators various questions, described below, that take into account multiple dimensions of how toxic

⁶We force beam search decoding to not use tokens from the prompt to prevent direct copying. Certain tokens appearing in the prompt such as punctuation are allowed.

⁷We compute the % of implicit samples using <https://github.com/RobertJGabriel/Google-profanity-words>, the same as ElSherief et al. (2021), also removing ambiguous terms (e.g., “bloody”).

machine-generated language presents a potential harm to readers. See Appendix B for an annotation screenshot and other study details.

Perceived hatefulness with respect to human or AI-authored text.

We first ask annotators to guess whether the statement’s author was a human or an AI system (HUMANORAI). Then, we ask whether the statement would be harmful to anyone if an AI system wrote it (HARMFULIFAI), as well as if a human wrote it (HARMFULIFHUMAN); we hypothesize that readers may have different standards for machine-generated text than human-written text. For all questions measuring harmfulness of text, we consider potential harm on a 1-5 scale with 1 being clearly benign and 5 indicating very offensive or abusive text.

Perceived intent of the writer. We ask readers whether statements were likely intended to be harmful (HARMFULINTENT), since some biased statements can be positively intended (e.g., benevolent sexism; Glick and Fiske, 1996). Additionally, we ask if the statement exhibits a positive stereotype (POSTEREO), which is also harmful (e.g., model minority myths; Cheryan and Bodenhausen, 2000).

Detailed harm explanations. To better understand how harm may be perpetrated against the minority group, we ask readers in-depth questions about text’s content, following Sap et al. (2020) and Olteanu et al. (2018). We ask whether or not the statement is lewd or sexual (LEWD), whether and how it references the targeted group or other groups (WHICHGROUP, GROUPFRAMING), whether it claims to be factual or opinion (FACTOROPINION).

4.2 Constructing TOXIGEN-HUMANVAL

Data and Setup. We selected 792 statements from TOXIGEN to include in our test set, such that no training statement had cosine similarity above 0.7 with any test statement. Each test statement was then rated by 3 annotators from a pool of 156 prequalified annotators from Amazon MTurk (See Appendix B for details).

Inter-annotator agreement. To investigate the quality of our annotations, we compute agreement on toxicity ratings.⁸ We find that annotators agreed moderately and are higher than or equal rates to prior work on hate speech annotation (Ross et al.,

⁸Specifically, we take the max of the HARMFULIFAI and HARMFULIFHUMAN scores and map it into three classes (scores <3: “non-toxic”, =3: “ambiguous”, >3: “toxic”).

Text	HUMANORAI (Predicted/True)	AVGSCORE (AI/Human)	INTENT	GROUPFRAMING
I started wearing a head scarf at the age of 12 because my mom told me that it was my duty as a woman to cover my hair	Human/AI (ALICE)	1.3/1.3	Benign	indirect reference, moral judgement
if the discussion is for the sake of the greater black community rather than for you, quit trying to hyper-centrally control it	Human/AI (top- k)	2.7/2.7	Unsure	direct reference, suggests solution, moral judgement
we do not need to mention trans people in an article about lgbtq+ issues	Human/AI (top- k)	4.0/4.3	Harm	direct reference, suggests solution
the mentally disabled should be able to live in their communities and not be forced into institutions.	Human/AI (ALICE)	1.0/1.0	Benign	direct reference, suggests solution, other framing

Table 3: Example responses from human evaluation where machine-generated text fools annotators into thinking the writer is human. Average toxicity scores are on a 1-5 scale (1 being benign and 5 being clearly offensive), and are averaged across annotator responses. We report scores for the case where annotators assume the writer/speaker is AI and the writer/speaker is human respectively.




Figure 4: Summary statistics for the human annotations on the evaluation set. Each statistic that the annotators are asked to evaluate is shown along the x-axis, while the y-axis gives the percentage of examples per annotated class (non-toxic, toxic, ambiguous).

2017; Sap et al., 2020), with a Fleiss’ $\kappa=0.46$ (Fleiss, 1971) and Krippendorff’s $\alpha=0.64$ (Krippendorff, 1980). In 55.17% of cases, all 3 annotators agree, while a majority ($\geq 2/3$) agree for 93.4%.

Human validation results. First, we find that our machine-generated statements are largely indistinguishable from human-written statements. For example—see Table 3—human annotators often




Figure 5: Avg. toxicity scores on a Likert scale of 1-5. Toxicity scores are similar across annotator-verified classes for a presumed AI speaker and human speaker.

predict that our text is generated by a human. In fact, on average 90.5% of machine-generated examples are thought to be human-written by a majority of annotators, as shown in Figure 4. We also note that harmful text confuses readers slightly more than non-harmful text: 92.9% of toxic examples are mislabeled as human-written compared to 90.2% for non-toxic. Most toxic examples are also hate speech (94.56%). While opinions are common in both toxic and non-toxic examples, most fact-claiming text is non-toxic.

Second, we find that demonstration-based prompting reliably generates toxic and benign statements about minority groups (§4.3). Further, for the machine-generated examples, we find that 30.2% are harmful (given a score of >3), while only 4% are ambiguous. This indicates that these data are sufficiently toxic or benign. We also find

that all identity groups covered by the dataset were represented in the human study (see Figure 3), and observe that the identity group referenced by the prompt is generally the same as the group referenced by the corresponding TOXIGEN text, though there is some deviation. This is likely due to GPT-3 conflating identities or mentioning multiple groups.

Interestingly, there is no significant difference in toxicity when we account for whether annotators perceive scores as written by humans or AI (Figure 5). This finding indicates that our machine-generated text is perceived as similarly harmful to human text. We also find that the most common framing tactic is “moral judgement”, or questioning the morality of an identity group, which has been linked to toxicity by prior work (Hoover et al., 2019).

4.3 Comparing Generation Methods

As further validation, we investigate whether ALICE-generated statements are more adversarial compared to top- k -generated ones. For 125 randomly-selected prompts (62 toxic and 63 non-toxic), we generate two statements: one with ALICE and one without (top- k). We then collect annotations for the 250 statements using the setup described in §4.1, and get toxicity scores from HateBERT.

We find that for top- k sampled sentences, the prompt label indeed matches the desired label (95.2% of non-toxic examples and 67.7% of toxic examples). For ALICE, 40.3% of toxic examples match the prompt label and 92.1% of non-toxic examples match. We also find that ALICE succeeds in fooling HateBERT (26.4% of ALICE-decoded sentences fool HateBERT vs. 16.8% of top- k sampled sentences). Finally, ALICE is effective for detoxifying generated text: the avg. human-annotated toxicity score for ALICE-decoded sentences with a toxic prompt is 2.97, compared to 3.75 for top- k . This difference is statistically significant with $p < 0.001$. ALICE therefore leads to harder, more ambiguous examples. We greatly expand on these findings in Appendix E with a larger scale human evaluation ($\sim 10,000$ samples) comparing sentences generated with and without ALICE.

5 Improving Toxicity Classifiers

To further showcase the usefulness of TOXIGEN, we investigate how it can enhance classifiers’ abilities to detect human-written and machine-generated implicit toxic language. We fine-tune

	Test Data	Finetune Data			
		None	ALICE	top- k	ALICE + top- k
HateBERT	SBF _{test}	0.60	0.66	0.65	0.71
	IHC	0.60	0.60	0.61	0.67
	DYNAHATE	0.47	0.54	0.59	0.66
	TOXIGEN-VAL	0.57	0.93	0.88	0.96
RoBERTa	SBF _{test}	0.65	0.70	0.67	0.70
	IHC	0.57	0.64	0.63	0.66
	DYNAHATE	0.49	0.51	0.50	0.54
	TOXIGEN-VAL	0.57	0.87	0.85	0.93

Table 4: AUC for HateBERT and RoBERTa both zero-shot and fine-tuned on 3 versions of our dataset: ALICE only, top- k only, and both combined. Since there are fewer ALICE samples than top- k , we downsample top- k for fair comparison via equal-sized datasets. ALICE + top- k combines these two datasets. Each model is evaluated on three external human-written datasets and the human-validated portion of TOXIGEN. Bolding denotes the best performance. In the zero-shot setting (first column) ALICE creates more challenging evaluation samples by attacking HateBERT and RoBERTa.

the widely-used HateBERT (Caselli et al., 2021) and ToxDectRoBERTa (Zhou et al., 2021) models on the training portion of TOXIGEN, using the prompt labels as proxies for a true toxicity label. Then, we compare the performance of the out-of-the-box models to those fine-tuned on TOXIGEN on three publicly available human-written datasets (IMPLICITHATECORPUS (ElSherief et al., 2021), the SOCIALBIASFRAMES test set (Sap et al., 2020), and DYNAHATE (Vidgen et al., 2021)) as well as the evaluation portion of our machine-generated dataset (TOXIGEN-HUMANVAL). To ablate the contribution of each decoding method, we also split TOXIGEN into equal numbers of ALICE-generated and top- k -generated examples.

Our results—see Table 4—show that fine-tuning HateBERT and ToxDectRoBERTa on TOXIGEN improves performance across all datasets. The improvement on human-written datasets shows that TOXIGEN can be used to improve existing classifiers, helping them better tackle the challenging human-generated implicit toxicity detection task. Fine-tuned HateBERT performs strongly on TOXIGEN-HUMANVAL, demonstrating that our data can successfully help guard against machine-generated toxicity.

6 Conclusions

In this work, we used a large language model to create and release TOXIGEN, a large-scale, balanced, and implicit toxic language dataset. TOXIGEN is

far larger than previous datasets, containing over 274k sentences, and is more diverse, including mentions of 13 minority groups at scale. The generated samples are balanced in terms of number of benign and toxic samples for each group. We proposed ALICE, an adversarial decoding scheme to evaluate robustness of toxicity classifiers and generate sentences to attack them, and showed the effectiveness of ALICE on a number of publicly-available toxicity detection systems. In our experiments, we showed that fine-tuning pre-trained hate classifiers on TOXIGEN can improve their performance on three popular *human*-generated toxicity datasets. We also conducted a human study on a subset of TOXIGEN, verifying that our generation methods successfully create challenging statements that annotators struggle to distinguish from human-written text: 90.5% of machine-generated examples were thought to be human-written.

7 Societal and Ethical Considerations

Risks in dataset release While the purpose of our work is to curate diverse and effective hate speech detection resources, our methods encourage a large language model to make its generation *more* toxic. This poses a potential misuse case where bad actors exploit these methods for nefarious purposes like spreading machine-generated hate speech. Still, ignoring this possibility does not make it go away and our work introduces an opportunity for the community to push back against harm towards minority groups. Our ultimate aim is to shift power dynamics to targets of oppression. Therefore, we do not consider identity dimensions that are historically the agents of oppression (e.g., whiteness, heterosexuality, able-bodied-ness). Please also note that there is still a lot that this dataset is not capturing about toxic language. Our annotations might not capture the full complexity of these issues related to human experiences. There is need for multi-disciplinary work to better understand these aspects.

ALICE The proposed method in this work attacks content filters via an adversarial game between two AI systems and thus passes the existing content filters—as we show for 5 publicly-available systems. It is important to leverage this and similar approaches to improve content filters and prevent large scale attacks against sensitive platforms.

Improving Toxicity Detection Effective classifiers for machine biases are required to combat the scale of online harm. Without such systems, minority groups are likely to be targeted by current (biased) systems. Our work is a significant step towards advancing this crucial classification task. Still, toxicity is inherently subjective (Sap et al., 2021). Therefore, moving beyond binary detection tasks to a focus on more nuanced labeling systems (ElSherief et al., 2021; Leonardelli et al., 2021) will prove crucial in developing responsible systems.

Relationship to Policy The topic of detecting and mitigating toxicity is relevant to the ongoing work and discussions in the space of policy and legislation for AI technology (Wischmeyer and Rademacher, 2020; Reich et al., 2021). Carefully crafted policy and regulation can play an important role in providing oversight into the development and deployment of content moderation systems and toxicity detection algorithms in practice (Benesch, 2020; Gillespie et al., 2020). Getting this right carries a crucial importance for the society as errors in content moderation can disproportionately affect minority groups (Sap et al., 2019). We see a path forward in which tools and techniques like those presented in this work are paired with human expertise and well-informed policy & regulation in bringing scalable and reliable solutions to practice. We acknowledge and encourage the critical role the NLP research community is poised to play in this inter-disciplinary effort.

Responsible AI Considerations Please also note that there is still a lot that this dataset is not capturing about what constitutes problematic language. Our annotations might not capture the full complexity of these issues, given problematic language is context-dependent, dynamic, and can manifest in different forms and different severities. Problematic language is also fundamentally a human-centric problem and should be studied in conjunction with human experience. There is need for multi-disciplinary work to better understand these aspects. Also note that this dataset only captures implicit toxicity (more precisely hate speech) for 13 identified minority groups, and due to its large scale can naturally be noisy. Our goal in this project is to provide the community with means to improve toxicity detection on implicit toxic language for the identified minority groups and there exists limitations to this dataset and models trained

on it which can potentially be the subject of future research, for example, including more target groups, a combination of them and so on that are not covered in our work.

8 Acknowledgements

We thank Azure AI Platform and Misha Bilenko for sponsoring this work and providing compute resources, Microsoft Research for supporting our large scale human study, and Alexandra Olteanu for her feedback on human evaluation. We also thank the crowdworkers for their time and effort.

References

- Peter Anderson, Basura Fernando, Mark Johnson, and Stephen Gould. 2017. Guided open vocabulary image captioning with constrained beam search. In *EMNLP*.
- Carolina Are. 2020. How instagram’s algorithm is censoring women and vulnerable users but helping online abusers. *Feminist media studies*, 20(5):741–744.
- Elizabeth Behm-Morawitz and Dana E Mastro. 2008. Mean girls? the influence of gender portrayals in teen movies on emerging adults’ Gender-Based attitudes and beliefs. *Journalism & mass communication quarterly*, 85(1):131–146.
- Susan Benesch. 2020. [Proposals for improved regulation of harmful online content](#). Technical report.
- Sid Black, Gao Leo, Phil Wang, Connor Leahy, and Stella Biderman. 2021. [GPT-Neo: Large Scale Autoregressive Language Modeling with Mesh-Tensorflow](#). If you use this software, please cite it using these metadata.
- Luke Breitfeller, Emily Ahn, David Jurgens, and Yulia Tsvetkov. 2019. Finding microaggressions in the wild: A case for locating elusive phenomena in social media posts. In *Proceedings of the 2019 Conference on Empirical Methods in Natural Language Processing and the 9th International Joint Conference on Natural Language Processing (EMNLP-IJCNLP)*, pages 1664–1674.
- Tom Brown, Benjamin Mann, Nick Ryder, Melanie Subbiah, Jared D Kaplan, Prafulla Dhariwal, Arvind Neelakantan, Pranav Shyam, Girish Sastry, Amanda Askell, Sandhini Agarwal, Ariel Herbert-Voss, Gretchen Krueger, Tom Henighan, Rewon Child, Aditya Ramesh, Daniel Ziegler, Jeffrey Wu, Clemens Winter, Chris Hesse, Mark Chen, Eric Sigler, Mateusz Litwin, Scott Gray, Benjamin Chess, Jack Clark, Christopher Berner, Sam McCandlish, Alec Radford, Ilya Sutskever, and Dario Amodei. 2020. [Language models are few-shot learners](#). In *Advances in Neural Information Processing Systems*, volume 33, pages 1877–1901.
- Tommaso Caselli, Valerio Basile, Jelena Mitrovic, and M. Granitzer. 2021. Hatebert: Retraining bert for abusive language detection in english. *ArXiv*, abs/2010.12472.
- S Cheryan and G V Bodenhausen. 2000. When positive stereotypes threaten intellectual performance: the psychological hazards of “model minority” status. *Psychological science*, 11(5):399–402.
- Elizabeth Clark, Tal August, Sofia Serrano, Nikita Haduong, Suchin Gururangan, and Noah A. Smith. 2021. [All that’s ‘human’ is not gold: Evaluating human evaluation of generated text](#). In *Proceedings of the 59th Annual Meeting of the Association for Computational Linguistics and the 11th International Joint Conference on Natural Language Processing (Volume 1: Long Papers)*, pages 7282–7296, Online. Association for Computational Linguistics.
- Aida Mostafazadeh Davani, Mark Díaz, and Vinodkumar Prabhakaran. 2022. [Dealing with Disagreements: Looking Beyond the Majority Vote in Subjective Annotations](#). *Transactions of the Association for Computational Linguistics*, 10:92–110.
- Thomas Davidson, Dana Warmesley, Michael Macy, and Ingmar Weber. 2017. Automated hate speech detection and the problem of offensive language. In *Proceedings of the International AAAI Conference on Web and Social Media*, volume 11.
- Thiago Dias Oliva, Dennys Marcelo Antonialli, and Alessandra Gomes. 2020. Fighting hate speech, silencing drag queens? artificial intelligence in content moderation and risks to LGBTQ voices online. *Sexuality & culture*.
- Ángel Díaz and Laura Hecht-Felella. 2021. Double standards in social media content moderation. Technical report, Brennan Center for Justice at New York University School of Law.
- Emily Dinan, Samuel Humeau, Bharath Chintagunta, and Jason Weston. 2019. [Build it break it fix it for dialogue safety: Robustness from adversarial human attack](#). In *Proceedings of the 2019 Conference on Empirical Methods in Natural Language Processing and the 9th International Joint Conference on Natural Language Processing (EMNLP-IJCNLP)*, pages 4537–4546, Hong Kong, China. Association for Computational Linguistics.
- Lucas Dixon, John Li, Jeffrey Sorensen, Nithum Thain, and Lucy Vasserman. 2018. Measuring and mitigating unintended bias in text classification. In *Proceedings of the 2018 AAAI/ACM Conference on AI, Ethics, and Society*, pages 67–73.
- Yao Dou, Maxwell Forbes, Rik Koncel-Kedziorski, Noah A Smith, and Yejin Choi. 2021. Scarecrow: A framework for scrutinizing machine text. *arXiv preprint arXiv:2107.01294*.

- Mai ElSherief, Caleb Ziems, David Muchlinski, Vaishnavi Anupindi, Jordyn Seybolt, Munmun De Choudhury, and Diyi Yang. 2021. Latent hatred: A benchmark for understanding implicit hate speech. *arXiv preprint arXiv:2109.05322*.
- Angela Fan, Mike Lewis, and Yann Dauphin. 2018. Hierarchical neural story generation. In *Proceedings of the 56th Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, pages 889–898.
- Joseph L. Fleiss. 1971. Measuring nominal scale agreement among many raters. *Psychological Bulletin*, 76:378–382.
- Antigoni Maria Founta, Constantinos Djouvas, Despoina Chatzakou, Ilias Leontiadis, Jeremy Blackburn, Gianluca Stringhini, Athena Vakali, Michael Sirivianos, and Nicolas Kourtellis. 2018. Large scale crowdsourcing and characterization of twitter abusive behavior. In *Twelfth International AAAI Conference on Web and Social Media*.
- Tianyu Gao, Adam Fisch, and Danqi Chen. 2021. Making pre-trained language models better few-shot learners. In *ACL*.
- Sam Gehman, Suchin Gururangan, Maarten Sap, Yejin Choi, and Noah A Smith. 2020. [Realtocixityprompts: Evaluating neural toxic degeneration in language models](#). In *Findings of EMNLP*.
- Ona de Gibert, Naiara Perez, Aitor García-Pablos, and Montse Cuadros. 2018. Hate speech dataset from a white supremacy forum. In *Proceedings of the 2nd Workshop on Abusive Language Online (ALW2)*, pages 11–20.
- Tarleton Gillespie, Patricia Aufderheide, Elinor Carmi, Ysabel Gerrard, Robert Gorwa, Ariadna Matamoros-Fernandez, Sarah T Roberts, Aram Sinnreich, and Sarah Myers West. 2020. [Expanding the debate about content moderation: Scholarly research agendas for the coming policy debates](#). *Internet Policy Review*, 9(4):Article number: 41–29.
- Peter Glick and Susan T Fiske. 1996. The ambivalent sexism inventory: Differentiating hostile and benevolent sexism. *Journal of personality and social psychology*, 70(3):491.
- Xiaochuang Han and Yulia Tsvetkov. 2020. [Fortifying toxic speech detectors against veiled toxicity](#). In *Proceedings of the 2020 Conference on Empirical Methods in Natural Language Processing (EMNLP)*, pages 7732–7739, Online. Association for Computational Linguistics.
- Chris Hokamp and Qun Liu. 2017. [Lexically constrained decoding for sequence generation using grid beam search](#). In *Proceedings of the 55th Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, pages 1535–1546, Vancouver, Canada. Association for Computational Linguistics.
- Ari Holtzman, Jan Buys, Maxwell Forbes, Antoine Bosselut, David Golub, and Yejin Choi. 2018. [Learning to write with cooperative discriminators](#). In *Proceedings of the 56th Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, pages 1638–1649, Melbourne, Australia. Association for Computational Linguistics.
- Joseph Hoover, Mohammad Atari, Aida M Davani, Brendan Kennedy, Gwenyth Portillo-Wightman, Leigh Yeh, Drew Kogon, and Morteza Dehghani. 2019. Bound in hatred: The role of group-based morality in acts of hate.
- David L Hudson, Jr. 2017. Hate speech online. <https://web.archive.org/web/20211115012316/https://www.freedomforuminstitute.org/first-amendment-center/topics/freedom-of-speech-2/internet-first-amendment/hate-speech-online/>. Accessed: 2021-11-14.
- Jonathan W Kanter, Monnica T Williams, Adam M Kuczynski, Katherine E Manbeck, Marlena Debreaux, and Daniel C Rosen. 2017. A preliminary report on the relationship between microaggressions against black people and racism among white college students. *Race and social problems*, 9(4):291–299.
- Brendan Kennedy, Mohammad Atari, Aida Mostafazadeh Davani, Leigh Yeh, Ali Omrani, Yehsong Kim, Kris Coombs, Shreya Havaladar, Gwenyth Portillo-Wightman, Elaine Gonzalez, et al. 2018. The gab hate corpus: A collection of 27k posts annotated for hate speech.
- Ben Krause, Akhilesh Deepak Gotmare, Bryan McCann, Nitish Shirish Keskar, Shafiq Joty, Richard Socher, and Nazneen Fatema Rajani. 2020. Gedi: Generative discriminator guided sequence generation. *arXiv preprint arXiv:2009.06367*.
- Klaus Krippendorff. 1980. Content analysis: an introduction to its methodology.
- Sumit Kumar and Raj Ratn Pranesh. 2021. Tweetblm: A hate speech dataset and analysis of black lives matter-related microblogs on twitter. *arXiv preprint arXiv:2108.12521*.
- Elisa Leonardelli, Stefano Menini, Alessio Palmero Aprosio, Marco Guerini, and Sara Tonelli. 2021. [Agreeing to disagree: Annotating offensive language datasets with annotators’ disagreement](#). In *Proceedings of the 2021 Conference on Empirical Methods in Natural Language Processing*, pages 10528–10539, Online and Punta Cana, Dominican Republic. Association for Computational Linguistics.
- Alisa Liu, Maarten Sap, Ximing Lu, Swabha Swayamdipta, Chandra Bhagavatula, Noah A. Smith, and Yejin Choi. 2021a. [Dexperts: Decoding-time controlled text generation with experts and anti-experts](#). In *ACL*.

- Jiachang Liu, Dinghan Shen, Yizhe Zhang, Bill Dolan, Lawrence Carin, and Weizhu Chen. 2021b. What makes good in-context examples for gpt-3? *arXiv preprint arXiv:2101.06804*.
- Pengfei Liu, Weizhe Yuan, Jinlan Fu, Zhengbao Jiang, Hiroaki Hayashi, and Graham Neubig. 2021c. Pre-train, prompt, and predict: A systematic survey of prompting methods in natural language processing. *arXiv preprint arXiv:2107.13586*.
- Ximing Lu, Peter West, Rowan Zellers, Ronan Le Bras, Chandra Bhagavatula, and Yejin Choi. 2021. **Neuro-Logic decoding: (un)supervised neural text generation with predicate logic constraints**. In *Proceedings of the 2021 Conference of the North American Chapter of the Association for Computational Linguistics: Human Language Technologies*, pages 4288–4299, Online. Association for Computational Linguistics.
- Sean MacAvaney, Hao-Ren Yao, Eugene Yang, Katina Russell, Nazli Goharian, and Ophir Frieder. 2019. Hate speech detection: Challenges and solutions. *PloS one*, 14(8):e0221152.
- Swaroop Mishra, Daniel Khoshabi, Chitta Baral, and Hannaneh Hajishirzi. 2021. Natural instructions: Benchmarking generalization to new tasks from natural language instructions. *arXiv preprint arXiv:2104.08773*.
- Kevin L Nadal. 2018. *Microaggressions and traumatic stress: Theory, research, and clinical treatment*. American Psychological Association.
- Kevin L Nadal, Katie E Griffin, Yinglee Wong, Sahran Hamit, and Morgan Rasmus. 2014. The impact of racial microaggressions on mental health: Counseling implications for clients of color. *Journal of counseling and development: JCD*, 92(1):57–66.
- Alexandra Olteanu, Carlos Castillo, Jeremy Boy, and Kush R. Varshney. 2018. The effect of extremist violence on hateful speech online. In *ICWSM*.
- Jeff Rasley, Samyam Rajbhandari, Olatunji Ruwase, and Yuxiong He. 2020. Deepspeed: System optimizations enable training deep learning models with over 100 billion parameters. In *Proceedings of the 26th ACM SIGKDD International Conference on Knowledge Discovery & Data Mining*, pages 3505–3506.
- Rob Reich, Mehran Sahami, and Jeremy M Weinstein. 2021. *System error: Where big tech went wrong and how we can reboot*. Hodder & Stoughton.
- Björn Ross, Michael Rist, Guillermo Carbonell, Benjamin Cabrera, Nils Kurowsky, and Michael Wojatzki. 2017. Measuring the reliability of hate speech annotations: The case of the european refugee crisis. *ArXiv*, abs/1701.08118.
- Paul Röttger, Bertie Vidgen, Dong Nguyen, Zeerak Waseem, Helen Margetts, and Janet Pierrehumbert. 2021. **HateCheck: Functional tests for hate speech detection models**. In *Proceedings of the 59th Annual Meeting of the Association for Computational Linguistics and the 11th International Joint Conference on Natural Language Processing (Volume 1: Long Papers)*, pages 41–58, Online. Association for Computational Linguistics.
- RWJF. 2017. Discrimination in america: experiences and views. <https://www.rwjf.org/en/library/research/2017/10/discrimination-in-america--experiences-and-views.html>. Accessed: 2019-11-5.
- Muniba Saleem and Craig A Anderson. 2013. Arabs as terrorists: Effects of stereotypes within violent contexts on attitudes, perceptions, and affect. *Psychology of violence*, 3(1):84–99.
- Maarten Sap, Dallas Card, Saadia Gabriel, Yejin Choi, and Noah A Smith. 2019. **The risk of racial bias in hate speech detection**. In *ACL*.
- Maarten Sap, Saadia Gabriel, Lianhui Qin, Dan Jurafsky, Noah A Smith, and Yejin Choi. 2020. Social bias frames: Reasoning about social and power implications of language. In *ACL*.
- Maarten Sap, Swabha Swayamdipta, Laura Vianna, Xuhui Zhou, Yejin Choi, and Noah A. Smith. 2021. **Annotators with attitudes: How annotator beliefs and identities bias toxic language detection**.
- Emily Sheng, Kai-Wei Chang, Premkumar Natarajan, and Nanyun Peng. 2019. **The woman worked as a babysitter: On biases in language generation**. In *Proceedings of the 2019 Conference on Empirical Methods in Natural Language Processing and the 9th International Joint Conference on Natural Language Processing (EMNLP-IJCNLP)*, pages 3407–3412, Hong Kong, China. Association for Computational Linguistics.
- Wiktor Soral, Michał Bilewicz, and Mikołaj Winiewski. 2018. Exposure to hate speech increases prejudice through desensitization. *Aggressive behavior*, 44(2):136–146.
- Derald Wing Sue, Christina M Capodilupo, Gina C Torino, Jennifer M Bucceri, Aisha M B Holder, Kevin L Nadal, and Marta Esquilin. 2007. Racial microaggressions in everyday life: implications for clinical practice. *The American psychologist*, 62(4):271–286.
- Bertie Vidgen, Tristan Thrush, Zeerak Waseem, and Douwe Kiela. 2021. **Learning from the worst: Dynamically generated datasets to improve online hate detection**. In *Proceedings of the 59th Annual Meeting of the Association for Computational Linguistics and the 11th International Joint Conference on Natural Language Processing (Volume 1: Long Papers)*, pages 1667–1682, Online. Association for Computational Linguistics.

- Ben Wang and Aran Komatsuzaki. 2021. GPT-J-6B: A 6 Billion Parameter Autoregressive Language Model. <https://github.com/kingoflolz/mesh-transformer-jax>.
- Zeerak Waseem. 2016. Are you a racist or am i seeing things? annotator influence on hate speech detection on twitter. In *Proceedings of the first workshop on NLP and computational social science*, pages 138–142.
- Michael Wiegand, Josef Ruppenhofer, and Elisabeth Eder. 2021. Implicitly abusive language—what does it actually look like and why are we not getting there? In *Proceedings of the 2021 Conference of the North American Chapter of the Association for Computational Linguistics: Human Language Technologies*, pages 576–587.
- Thomas Wischmeyer and Timo Rademacher, editors. 2020. *Regulating Artificial Intelligence*. Springer, Cham.
- Albert Xu, Eshaan Pathak, Eric Wallace, Suchin Gururangan, Maarten Sap, and Dan Klein. 2021. *Detoxifying language models risks marginalizing minority voices*. In *NAACL*.
- Kevin Yang and Dan Klein. 2021. *FUDGE: Controlled text generation with future discriminators*. In *Proceedings of the 2021 Conference of the North American Chapter of the Association for Computational Linguistics: Human Language Technologies*, pages 3511–3535, Online. Association for Computational Linguistics.
- Danyaal Yasin. 2018. Black and banned: Who is free speech for? <https://www.indexoncensorship.org/2018/09/black-and-banned-who-is-free-speech-for/>. Accessed: 2018-12-6.
- Marcos Zampieri, S. Malmasi, Preslav Nakov, Sara Rosenthal, N. Farra, and Ritesh Kumar. 2019. Predicting the type and target of offensive posts in social media. In *NAACL*.
- Xuhui Zhou, Maarten Sap, Swabha Swayamdipta, Yejin Choi, and Noah A. Smith. 2021. Challenges in automated debiasing for toxic language detection. In *EACL*.

Supplementary Materials

A Generation Details

To generate sentences for a given minority group, we sample 5 random sentences from the corresponding set of examples, then join them into one string with each example being preceded by a hyphen (“-”) and ending with a newline character (“\n”). By appending an extra hyphen to the end of the prompt, LLMs write a new sentence matching the style of the presented examples. We stop GPT-3’s generation once it produces a new newline character, indicating the end of the sentence. For each generated sentence, we use a new, randomly-selected set of 5 random sentences.

A.1 Language Model Selection

While we use GPT-3 to generate statements in this work, in principle, our methods can be used with any models that generate realistic text, such as GPT-Neo (Black et al., 2021), GPT-J (Wang and Komatsuzaki, 2021), or Turing-NLG (Rasley et al., 2020)

B Human Validation Details

B.1 Selecting MTurk Workers

For human validation, we select 156 MTurk workers with prior experience annotating toxic language (Sap et al., 2020). 51 of these workers participated in data annotation. We collect worker demographics using an optional survey at the end of the annotation task. We find that 56.9% identify as White, 9.8% as Black, 3.9% as Hispanic, 3.9% as Asian and 5.9% as Other. Also, 45.1% of workers identify as female, 37.3% as male and 2% as non-binary. The majority of workers are between 25 and 45 (58.8%). Politically, 25.5% of workers identify as left-leaning, 23.5% as very left-leaning, 13.7% as moderate, 17.6% as right-leaning and 3.9% as very right-leaning.⁹ Lastly, we find that 5.9% of workers also identify as LGBTQ+ and 2% identify as Pacific Islander.

B.2 Annotation Interface

Figure 6 shows a screenshot of the annotation interface given to the Amazon Mechanical Turk workers. Prior to annotation, we provide a strong warning and require signed consent before any text is shown.

⁹The remaining workers chose not to respond for these questions.

C How does perplexity change across groups?

Our decoding approaches should ideally generate low-perplexity sentences. We measure the perplexity assigned by a pre-trained language model across different minority groups for sentences generated with and without ALICE. This will give us an idea of how good the set of sentences are from the perspective of the pre-trained language model in terms of perplexity. We use GPT-2 model from Huggingface to measure perplexity. As some sentences have extremely high perplexity according to GPT-2, we drop sentences (roughly 10% of the dataset) with perplexity over 500 for this analysis. As shown in Table 5, the ALICE-generated sentences have significantly lower perplexity than top- k across all minority groups. We also find that the average perplexity can range significantly between subgroups, though perplexity varies more for top- k -generated text. Interestingly, text mentioning Black people is deemed most-likely across the board, while the least-likely generations differ by generation method: amongst the ALICE-generated text, sentences mentioning Latino people is the least likely, while for top- k , text mentioning Women is the least likely. In all cases, ALICE generates text with up to 5 times lower perplexity than regular decoding.

Group	ALICE	top- k
Black	16.10	86.88
Asian	17.75	108.83
Native Am.	25.92	103.87
Muslim	17.16	84.92
Latino	36.69	96.68
Jewish	19.37	96.71
Chinese	33.60	121.54
LGBTQ+	18.15	87.93
Mental Dis.	21.22	92.21
Physical Dis.	30.46	129.15
Mexican	28.36	113.62
Women	21.44	131.52
Middle Eastern	30.71	127.95
Total	23.54	105.31

Table 5: Perplexity for different minority groups. Sentences with perplexity over 500 are dropped.

D Does generated text actually mention the targeted groups?

In the human validation study (§4), we ask annotators to determine whether or not the text actually includes references to the targeted groups;

Group	ALICE	top- <i>k</i>
Black	.87	.83
Asian	.62	.71
Native Am.	.96	.73
Latino	1.0	.72
Jewish	.60	.67
Muslim	.96	.89
Chinese	.73	.86
Mexican	.84	.91
Middle Eastern	.81	.77
LGBTQ+	.91	.97
Women	.97	.90
Mental Dis.	.84	.78
Physical Dis.	.86	.78
All groups	.84	.81

Table 6: Proportion of generated sentences that mention targeted identity groups in text generated with and without ALICE.

each prompt was generated with one group in mind. Here, we compare the proportion of text that mentions each group, split by decoding method. As shown in Table 6, we find that both ALICE and top-*k* generate text that mentions corresponding minority group in the prompt almost equally good (slightly better for ALICE), though the exact proportion changes by the group. For instance, in text generated for Latino people, ALICE has a 100% hit rate, while top-*k* has only 72%. However, for text mention LGBTQ+ people, top-*k* text succeeds to mention them 97% of the time while ALICE has only 91%. These values may depend on the underlying language model: in our case, GPT-3 may have been trained on less Latino-mentioning text and therefore benefit more from controlled decoding.

E Analysis of Large-Scale Human Validation

Summary Statistics. In addition to the human-validated evaluation set described in Section 4, we obtain labels for 8,960 randomly sampled training examples using the same annotation framework and pool of MTurk workers. This sample is evenly split between top-*k* and ALICE generated texts (50.9% for top-*k*, 49.1% for ALICE). Please note that the samples are drawn randomly from TOX-IGEN training data and we did not enforce having the same prompt for top-*k* and ALICE. The analysis for having the same prompt for top-*k* and ALICE has already been done in §4.3. In Figure 7, we show that average toxicity scores are similar

for the two decoding methods given a prompt label, though ALICE-generated texts have a slightly higher average toxicity for generations from toxic-labeled prompts and a lower average toxicity for generations from non-toxic prompts than top-*k*. We observe that 66.86% of ALICE-generated texts with a toxic prompt label are actually toxic (compared to 57.91% of top-*k* examples) and 93.21% of ALICE-generated texts with a non-toxic prompt label are actually non-toxic (compared to 90.01% of top-*k* examples). We also find that *ALICE is more effective at generating adversarial language* - 58.97% of toxic ALICE-generated examples fool HateBERT, compared to 26.88% of toxic top-*k* generated examples. ALICE-generated non-toxic examples also fool HateBERT more often than top-*k*, though the difference is smaller (15.51% of ALICE-generated non-toxic examples vs. 11.35% of top-*k* generations). Lastly, *ALICE-generated examples lead to more consistent target group references*. At least one annotator identified a direct or indirect reference to the exact target group for 70.4% of top-*k* generated examples compared to 78.3% of ALICE-generated examples. A full breakdown is given in Figure 8.

Subjectivity in Annotation. As we address broadly in Section 7, subjectivity is an area of concern for annotation of toxicity. Prior work has pointed out the role that annotators’ belief systems and sociodemographic backgrounds play in their perception of toxicity (Sap et al., 2019, 2021; Davani et al., 2022). This is highlighted by such training examples as

“I’d like to thank you.. I don’t think I’d have made it if I didn’t find MGTOW!
You guys are great!”

Annotators predicted this example to be non-toxic, likely due to not recognizing MGTOW as a misogynistic group.

F Example Prompt

Prompt engineering can have significant effects on the quality of text generated by language models. Following the lead of other recent works, we use demonstration-based prompting, and introduce demonstrations to encourage language models to generate group-mentioning text. For example, a prompt that generates neutral mentions of Native American people is:

Test Data		Finetune Data			
		None	ALICE	top- <i>k</i>	ALICE + top- <i>k</i>
HB	TOXIGEN-VAL ALICE subset	0.44	1.00	0.80	0.99
	TOXIGEN-VAL top- <i>k</i> subset	0.72	0.80	0.95	0.92
RB	TOXIGEN-VAL ALICE subset	0.59	0.92	0.81	0.93
	TOXIGEN-VAL top- <i>k</i> subset	0.65	0.77	0.89	0.90

Table 7: Breaking the TOXIGEN-VAL test set into subsets with and without ALICE. HB denotes HateBERT, RB is ToxDectRoBERTa.

“

- it is a common misconception that native americans are all savages but they have a rich culture and history\n
- indigenous people are often left out of conversations around policies that impact them\n
- some native american communities have reclaimed the term indian, despite a complex history\n
- some native americans were granted land after the british were defeated in the american revolution by arrows\n
- indigenous people have lived in north america for more than 12,000 years\n
- there was a lot of resentment in the 19th century to native americans and some indians were even killed off\n

-
”

The trailing hyphen encourages the language model to provide another sentence in the growing list that resembles the prior examples, ending in \n. All of our generated prompts (26,000) are released with the dataset.

G Releasing a Pretrained Model and its Propagated Labels

We further finetune and release a RoBERTa classifier on the 8,960 human-annotated sampled in

TOXIGEN, beginning with the weights from (Zhou et al., 2021). Along with our publicly-available code, this pretrained model will serve as an entry point for community engagement with our work. We run this pretrained model on the full TOXIGEN dataset, collecting its predictions and release them along with TOXIGEN. These new labels may serve to correct some mislabeling.

H Dataset Description

We release TOXIGEN as a dataframe with the following fields: **prompt** contains the prompts we use for each generation. **generation** is the TOXIGEN generated text. **generation method** denotes whether or not ALICE was used to generate the corresponding generation. If this value is ALICE, then ALICE was used, if it is top-*k*, then ALICE was not used. **prompt_label** is the binary value indicating whether or not the prompt is toxic (1 is toxic, 0 is benign), and therefore the generation should be toxic as well. This label is slightly noisy, though largely accurate—as deemed by human annotators. **group** indicates for which group the prompt was generated. Finally, **roberta_prediction** is the probability predicted by our corresponding RoBERTa model for each instance. This field can be used as propagated labels according to this model.

I Further comparing toxicity classifiers

We also compare finetuning classifiers on subsets of TOXIGEN-VAL with and without ALICE, shown in Table 7. As expected, when finetuning on each subset individually, performance is strong on their respective evaluation sets. Further, without any finetuning, each model performs worse on the ALICE-generated data, indicating ALICE successfully generates data that are more confusing to each model.

Instructions

Read a text (potentially generated by an AI system) and tell us whether or not it contains harmful content (this should take less than 3 minutes). Note: this is a pilot task, more HITS will be available in the future.

Task

Text: They have been bred to be good at sports and entertainment, but not much else. Targeted Group: Black folks

☐ I'm having a lot of trouble understanding this text.

Answer questions about the text.

1) Do you think the text is human- or machine-generated?

☐ Human ☐ AI

2) In your opinion, would this text be seen as harmful to anyone **if an AI system wrote it**?

Not at all ☐ ☐ ☐ ☐ Very much so

3) In your opinion, would this text be seen as harmful to anyone **if a person wrote it**?

Not at all ☐ ☐ ☐ ☐ Very much so

4) In your opinion, does the text express positive stereotyping?

E.g. "all asians are good at math"

☐ Yes ☐ Maybe ☐ No

5) In your opinion, was the text intentionally written to be harmful to anyone?

E.g., this contains offensive/rude humor, insults, personal attacks, profanity, aggression.

Not at all ☐ ☐ ☐ ☐ Very much so

6) In your opinion, does this text contain or allude to sexual content/acts?

E.g. euphemism, double entendre, etc.

☐ Yes ☐ Maybe ☐ No

7) In your opinion, does the text refer to the given targeted group?

Please select one option from the list below.

-- please select --

If the text refers directly/indirectly to one or more other groups, please list them below:
(E.g., racial, gender identity, sexuality, disability, immigrant, veteran, etc.); use commas to separate groups.

e.g. women

8) In your opinion, how does the text refer to the targeted individual/group?

Please select one option from the list below.

-- please select --

If you selected "Other," please provide a single sentence explaining how the individual/group is being targeted in general terms. If no individual or group is targeted, you can leave this blank:

e.g. positive stereotyping

9) In your opinion, does the text explicitly claim to be factual?

Please select one option from the list below. If the text is explicitly factual, it should present content in a way that indicates an intent to inform, providing (possibly falsified) information like demographic-related statistics. If the text is explicitly opinion, it should be stated in the text that the content is not fact (e.g. "these are just my thoughts, but...").

-- please select --

Figure 6: Annotation setup for evaluating offensiveness of GPT-3 generations.




Figure 7: Average human-validated toxicity scores for training set examples based on prompt label (toxic vs. non-toxic) and decoding method (top-*k* vs. ALICE).




Figure 8: Comparing the proportion of identity group mentions that were desired based on the *prompts* vs. that were *generated*, in our large-scale validated training set. We include the actual proportions as data labels.